

Introduction to Data Science

Daniel Gutierrez, Data Scientist
Los Angeles, Calif.

Course Outcomes

- Provide an overview of Data Science
- Discuss the *Data Science Venn Diagram*
- Define the *Data Science Process*.
- Present an *Introduction to R Programming Part I*
- Start filling your data science toolbox with useful R programming techniques

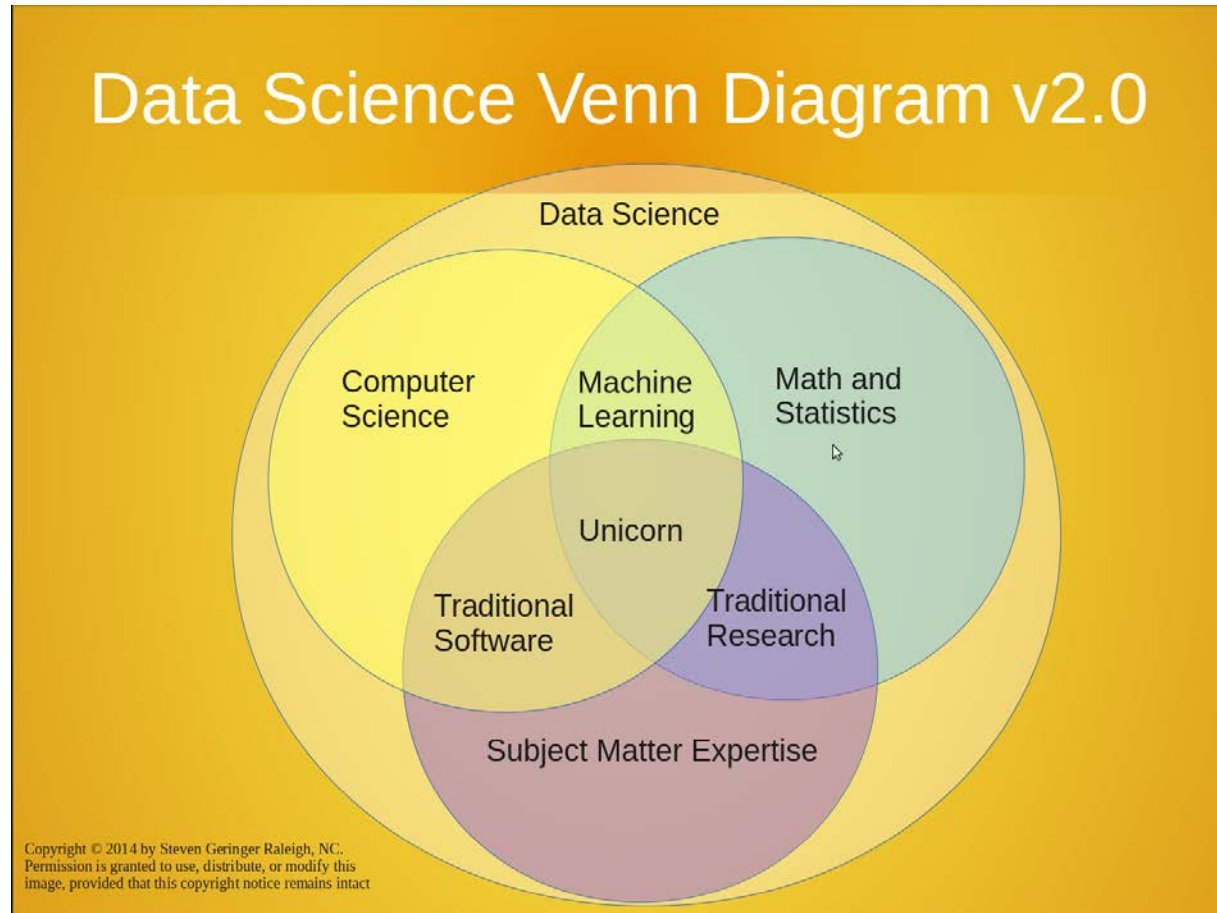
Lesson Objectives

- Fully understand the *Data Science Process*
- Provide a brief history of the R statistical environment
- Discuss installing, configuring and using R and Rstudio
- Talk about writing R scripts using basic language constructs and data types
- Understand the atomic classes in R
- Show how to code assignment statements
- Define a number of useful R objects: vectors, lists, matrices, factors, data frames, arrays
- Show how to create sequences
- Review object attributes such as names and dimensionality
- Discuss the importance of commenting your R code
- Give examples of coercion
- Show how to handle missing values: NA (not available), NaN (not a number) and NULL

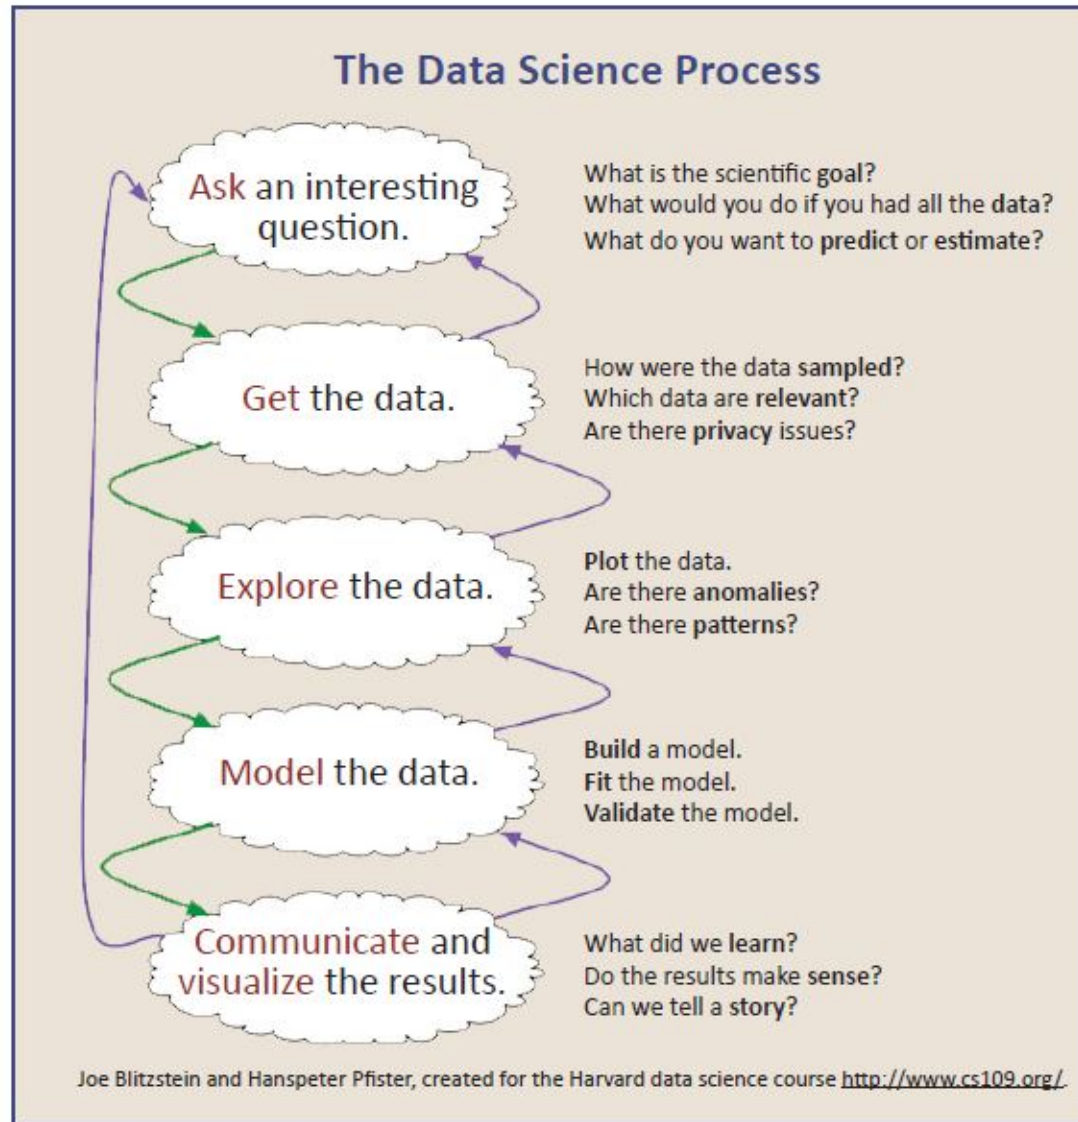
Data Science vs. Data Engineering

- A big confusion exists in our industry. What are the data scientist skillsets?
- Many companies advertise for “data scientist” but expect a vast array of skills
- Many companies actually seek a “unicorn”
- Realistically, these companies need a data science “team” consisting of both data scientists plus data engineers
- This class is for data scientists

The Data Science Venn Diagram



Data Science Process



Not So Brief History of R

- What is R?
- Historically, R is a dialect of the S language
- S is a language that was originally developed by John Chambers and others at Bell labs
- S was initiated in 1976 as an internal statistical analysis environment – originally implemented as Fortran libraries
- Early versions of the language did not contain functions for statistical modeling
- In 1988 the system was rewritten in C and began to resemble the system that we have today (this was Version 3 of the language). The book *Statistical Models in S* by Chambers and Hastie (the white book) documents the statistical analysis functionality.

Not so Brief History of R (continued)

- Version 4 of the S language was released in 1998. The book *Programming with Data* by John Chambers (the green book) documents this version of the language
- In 1993 Bell Labs gave StatSci (now Insightful Corp.) an exclusive license to develop and sell the S language
- In 2004 Insightful purchased the S language from Lucent for \$2 million
- Insightful sells its implementation of the S language under the product name S-PLUS
- In 2008 Insightful is acquired by TIBCO for \$25 million (TIBCO is a prominent player in the R market today)

Not so Brief History of R (continued)

- 1991: Created in New Zealand by Ross Ihaka and Robert Gentleman. Their experience developing R is documented in a 1996 JCGS paper
- 1993: First announcement of R to the public
- 1995: Martin Mächler convinces Ross and Robert to use the GNU General Public License to make R free software
- 1997: The R Core Group is formed (containing some people associated with S-PLUS). The core group controls the source code for R
- 2000: R Version 1.0.0 is released
- 2019: R version 3.5.3 is released on March 11, 2019

Not so Brief History of R (continued)

- Quite lean, as far as software goes; functionality is divided into modular “packages”
- Graphics capabilities very sophisticated and better than most stat packages.
- Useful for interactive work, but contains a powerful programming language for developing new tools (user → programmer)
- Very active and vibrant user community; R-bloggers (www.r-bloggers.com) and Stack Overflow (www.stackoverflow.com)

The R Environment

- The R system is divided into 2 conceptual parts:
 1. The “base” R system that you download from the CRAN website
 2. Everything else!
- The “everything else” consists of many other packages:
 - There are more than 13,914 packages on CRAN that have been developed by users and programmers around the world.
 - People often make special purpose packages available on their personal websites; there is no reliable way to keep track of how many packages are available in this fashion.

The R Environment

- In order to use an R package you must first determine if it exists
 - Find the package via Google e.g. “time series in R”
 - Use CRAN page such as: <http://cran.r-project.org/web/packages/timeSeries/index.html>
 - Download “Reference Manual” and vignettes (if any)
- ```
> install.packages("timeSeries")
> library(timeSeries)
```

# The R Environment

- Available from CRAN (<http://cran.r-project.org>)
  - The Comprehensive R Archive Network
  - An introduction to R
  - Writing R packages
  - R data import/export
  - R installation and administration (mostly for building R from source code)
  - R internals (not for the faint of heart)

# Installing R

- Download R from CRAN (<http://cran.r-project.org>)
- Select your environment: Windows, Mac, or Linux
- Click on “Install R for the first time”
- Click on “Download R 3.5.3 for Windows” for example
- Follow the useful installation prompts
- NOTE: you must install R before RStudio

# Installing RStudio

- Download the RStudio IDE from <http://www.rstudio.com/>
- Click on “Download RStudio” button
- Click on “RStudio Desktop”
- Click on “Download RStudio Desktop” button
- Click on the installer that matches your system
- Follow the useful installation prompts
- You should now see both the R and RStudio icons on your desktop

# R Fundamentals – Part 1

The screenshot displays the RStudio interface with the following components:

- Script Editor:** Contains a file named 'MODULE1 R Fundamentals.R' with the following code:

```
1 # UC Davis Extension - Introduction to Data Science
2 #
3 # MODULE 1 - R fundamentals
4 #
5 #
6 # (c) Copyright 2015 - AMULET Analytics
7 # -----
8
9
10 # Installing new R updates
11
12 install.packages("installr")
13
14 library(installr)
15
16 R.version.string
17
18 updateR() # updating R.
```
- Console:** Shows the R version and platform information:

```
R version 3.2.1 (2015-06-18) -- "world-famous Astronaut"
Copyright (C) 2015 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> x <- factor(c("yes", "yes", "no", "yes", "no"))
```
- Environment Pane:** Displays the 'Global Environment' with variables:
  - Data:** 'm' is an integer vector of length 2 with values 1 and 2.
  - Values:** 'x' is an integer vector of length 10 with values 1 through 10. 'y' is a factor with 3 levels: 'a', 'b', and 'c', with counts of 2, 1, and 3 respectively.
- Package List:** A table of installed and available packages:

| Name       | Description                                                      | Version   |
|------------|------------------------------------------------------------------|-----------|
| alr4       | Data to accompany Applied Linear Regression 4rd edition          | 1.0.5     |
| BB         | Solving and Optimizing Large-Scale Nonlinear Systems             | 2014.10-1 |
| car        | Companion to Applied Regression                                  | 2.0-25    |
| colorspace | Color Space Manipulation                                         | 1.2-6     |
| crayon     | Colored Terminal Output                                          | 1.3.0     |
| dfoptim    | Derivative-free Optimization                                     | 2011.8-1  |
| dichromat  | Color Schemes for Dichromats                                     | 2.0-0     |
| digest     | Create Cryptographic Hash Digests of R Objects                   | 0.6.8     |
| e1071      | Misc Functions of the Department of Statistics (e1071), TU Wien  | 1.6-4     |
| effects    | Effect Displays for Linear, Generalized Linear, and Other Models | 3.0-4     |
| evaluate   | Parsing and Evaluation Tools that Provide More Details than      | 0.7       |



# R Fundamentals – Part 1

- **Tips on Developing with R**
- On your local machine, set up a distinct folder (directory) for each R project. Don't mix R projects in the same folder
- Use `Session -> Set Working Directory` in RStudio to point to your project folder
- Use `File -> New File -> R Script` to open up a new R script in the code editor
- Start coding and using the R command line
- Use `Session -> Save Workspace As` to create an environment file containing all your variables. This will be named with the name you provide plus an `.RDATA` extension
- Next time, start up Rstudio with the `.RDATA` file to resume

# R Fundamentals – Part 1

- R has five basic or “atomic” classes of objects:
  - Character
  - Numeric (real numbers)
  - Integer
  - Complex
  - Logical (True / False)
- The most basic object is a vector
  - A vector can only contain objects of the same class
  - BUT: The one exception is a *list*, which is represented as a vector but can contain objects of different classes (indeed, that’s usually why we use them)

# R Fundamentals – Part 1

- At the R command prompt we generally type what are called *expressions*. You type the expression and R evaluates it.
- The `<-` symbol is the assignment operator (very similar in function to `=` assignment operator found in other programming environments) which assigns what's on the right to what's on the left
- The `#` character indicates a *comment*. Anything to the right of the `#` (including the `#` itself) is ignored as part of the R expression. Commenting your R scripts is highly recommended!

# R Fundamentals – Part 1

- Numbers in R are generally treated as numeric objects (i.e., double precision real numbers)
- If you explicitly want an integer, you need to specify the L suffix
- Ex: Entering 1 gives you a numeric object; entering 1L explicitly gives you an integer.
- There is also a special number `Inf` which represents infinity; e.g. `1 / 0`; `Inf` can be used in ordinary calculations; e.g. `1 / Inf` is 0
- The value `NaN` represents an undefined value (“not a number”); e.g. `0 / 0`; `NaN` can also be thought of as a missing value (more on that later)

# R Fundamentals – Part 1

- In addition to the atomic types, here are some additional object types available in R:
  - Vectors
  - Lists
  - Matrices
  - Factors
  - Data frames
  - Arrays

# R Fundamentals – Part 1

- R objects also can have attributes
  - Names
  - dimnames
  - Dimensions (e.g., matrices, arrays)
  - Class
  - Length
  - Other user-defined attributes/metadata

# R Fundamentals – Part 1

- Matrices are vectors with a *dimension* attribute. The dimension attribute is itself an integer vector of length 2 (nrow, ncol)
- Matrices are constructed *column-wise*, so entries can be thought of starting in the “upper left” corner and running down the columns.
- Lists are a special type of vector that can contain elements of different classes. The List is a very important data type in R and you should get to know them well.

# R Fundamentals – Part 1

- Factors are used to represent categorical data. Factors can be unordered or ordered. One can think of a factor as an integer vector where each integer has a *label*.
  - Factors are treated specifically by modelling functions like `lm()` and `glm()`
  - Using factors with labels is *better* than using integers because factors are self-describing; having a variable that has values “Male” and “Female” is better than a variable that has values 1 and 2.



# R Fundamentals – Part 1

- Missing values are denoted by NA or NaN for undefined mathematical operations (e.g. division by zero).
  - `is.na()` is used to test objects if they are NA
  - `is.nan()` is used to test for NaN
  - NA values have a class also, so there are integer NA, character NA, etc.
  - A NaN value is also NA but the converse is not true

# R Fundamentals – Part 1

- Data frames are used to store tabular data
  - They are represented as a special type of list where every element of the list has to have the same length
  - Each element of the list can be thought of as a column and the length of each element of the list is the number of rows
  - Unlike matrices, data frames can store different classes of objects in each column (just like lists); matrices must have every element be the same class
  - Data frames also have a special attribute called `row.names`
  - Data frames are usually created by calling `read.table()` or `read.csv()`
  - Can be converted to a matrix by calling `data.matrix()`

# R Fundamentals – Part 1

- Our goal for course Weeks 1 through 4 is to present a collection of R fundamentals and start filling your data science toolbox with a collection of programming techniques

# Code Modules

- WEEK 1-1 Code Module – updating R
- WEEK 1-2 Code Module – vectors
- WEEK 1-3 Code Module – coercion
- WEEK 1-4 Code Module – matrices
- WEEK 1-5 Code Module – lists
- WEEK 1-6 Code Module – factors
- WEEK 1-7 Code Module – missing values
- WEEK 1-8 Code Module – data frames
- WEEK 1-9 Code Module – names
- WEEK 1-10 Code Module - arrays

# Summary

- In WEEK 1 of Introduction to Data Science, we presented all of the introductory material for the course including the Data Science Venn diagram, and the Data Science Process.
- We also introduced R by giving a short history of the language as well as instructions for installing both R and RStudio.
- Next, we start our journey through the R language by learning some basics.